



(A) Mixed audio spectrogram — both speakers visible across full spectrum

Whisper transcribes correctly: CER = 0.44



(B) Oracle-separated Speaker 1 — speech (0–2.1s) then SILENCE

Whisper transcribes OK: CER = 0.44 (trailing silence is less harmful)



(C) Oracle-separated Speaker 2 — LEADING SILENCE (0–2.0s) then speech

Blank spectrogram triggers compression-seeking attractor: CER = 24.25, CR = 16.33

